



Object-Oriented Programming, Functional Programming and R

Author(s): John M. Chambers

Source: *Statistical Science*, May 2014, Vol. 29, No. 2 (May 2014), pp. 167-180

Published by: Institute of Mathematical Statistics

Stable URL: <https://www.jstor.org/stable/43288468>

JSTOR is a not-for-profit service that helps scholars, researchers, and students discover, use, and build upon a wide range of content in a trusted digital archive. We use information technology and tools to increase productivity and facilitate new forms of scholarship. For more information about JSTOR, please contact support@jstor.org.

Your use of the JSTOR archive indicates your acceptance of the Terms & Conditions of Use, available at <https://about.jstor.org/terms>



JSTOR

Institute of Mathematical Statistics is collaborating with JSTOR to digitize, preserve and extend access to *Statistical Science*

Object-Oriented Programming, Functional Programming and R

John M. Chambers

Abstract. This paper reviews some programming techniques in R that have proved useful, particularly for substantial projects. These include several versions of object-oriented programming, used in a large number of R packages. The review tries to clarify the origins and ideas behind the various versions, each of which is valuable in the appropriate context.

R has also been strongly influenced by the ideas of functional programming and, in particular, by the desire to combine functional with object oriented programming.

To clarify how this particular mix of ideas has turned out in the current R language and supporting software, the paper will first review the basic ideas behind object-oriented and functional programming, and then examine the evolution of R with these ideas providing context.

Functional programming supports well-defined, defensible software giving reproducible results. Object-oriented programming is the mechanism *par excellence* for managing complexity while keeping things simple for the user. The two paradigms have been valuable in supporting major software for fitting models to data and numerous other statistical applications.

The paradigms have been adopted, and adapted, distinctively in R. Functional programming motivates much of R but R does not enforce the paradigm. Object-oriented programming from a functional perspective differs from that used in non-functional languages, a distinction that needs to be emphasized to avoid confusion.

R initially replicated the S language from Bell Labs, which in turn was strongly influenced by earlier program libraries. At each stage, new ideas have been added, but the previous software continues to show its influence in the design as well. Outlining the evolution will further clarify why we currently have this somewhat unusual combination of ideas.

Key words and phrases: Programming languages, functional programming, object-oriented programming.

1. INTRODUCTION

R has become an important medium for communicating new methodology in statistics and related technology. References to the supporting R software frequently accompany journal articles or other publications describing new results. The software is available to other R users, ideally as a package in a standard

repository. The benefits for statistics as a discipline are considerable: The community has rapid access to new ideas in a free, open-source format as software that can in most cases be installed and used immediately by those interested in the statistical techniques. The user community has both created and benefited from this resource.

This paper examines two of the most significant paradigms in programming languages generally: object-oriented programming (OOP) and functional programming. R makes use of both, but in its own

John M. Chambers is Consulting Professor, Department of Statistics, Stanford University, Stanford, California 94305-4065, USA (e-mail: jmc@stat.stanford.edu).

way. Both paradigms are valuable for serious programming with the language. But in both cases, understanding the relevant ideas in the context of R is needed to avoid confusion. The confusion sometimes arises, in both cases, from applying to R interpretations of the paradigms that apply to other languages but not to this one. Section 2 of the paper will review the ideas, generally and in their R versions, with the goal of clarifying the basics. Given the importance of R software to the community, creators of new R software should benefit from understanding these concepts.

We will also examine in Section 3 of the paper the evolution that led to these versions of functional programming and OOP. The prime motivation was not language design in the abstract but to provide the tools needed for research and data analysis by the user community at the time. R originally reproduced the functionality of the S language at Bell Labs, which itself had evolved through several stages beginning in the late 1970s and which was in turn based on earlier statistical software libraries, mainly in Fortran.

R added important new ideas and has continued to evolve, but the main contents inherited through S shaped the capabilities and the approach to statistical computing. In a surprising number of areas, what we think of as “the R way” of organizing the computations actually reflects software developed twenty years or more before R existed.

Having been involved in all the stages, I am naturally inclined to a historical perspective, but it is also the case that the history itself had substantial impact on the results. It may be comforting to view programming languages as abstract definitions, but in practice they evolve from the needs, interests and limitations of their creators and users.

2. FUNCTIONAL AND OBJECT-ORIENTED PROGRAMMING: THE MAIN IDEAS

Functional and object-oriented programming fit naturally into statistical applications and into R. The original motivating use case, fitting models to data, remains compelling. An expression such as

```
irisFit <- lm(Sepal.Width ~
             . - Sepal.Length, iris)
```

calls a function that creates an object representing the linear model specified by the first argument, applied to the data specified by the second argument. The computation is functional, well-defined by the arguments. It returns an object whose properties provide the information needed to study and work with the fitted model.

Other functions and other objects can adapt to different models in a form that is convenient for both the user and the implementer.

Principles of functional programming guide us in writing reliable, reproducible functions for the different models. Object-oriented programming provides tools for defining the model objects clearly, and adapting to new ideas and new forms of models. Section 3.4 goes into details of the R implementations.

As they have been realized in R, both paradigms center on a few, intuitive concepts. The details are more complicated, as they usually are. In the case of functional programming, the realization in R is only partial, reflecting the language’s origins as well as practical considerations. In the case of OOP, there are now at least three realizations of the ideas in R, using two different paradigms. All three have significant applications and practical value.

Despite all these devilish details, the main ideas remain visible and useful, particularly when programming serious applications using the language.

2.1 Functional Programming

For our purposes, the main principles of functional programming can be summarized as follows:

1. Programming consists largely of defining *functions*.
2. A function definition in the language, like a function in mathematics, implies that a function call returns a unique value corresponding to each valid set of arguments, but *only* dependent on these arguments.
3. A function call has no side effects that could alter other computations.

The implication of the second point is that functions in the programming language are mappings from the allowed set of arguments to some range of output values. In particular, the returned value should not depend on other quantities that affect the “state” of the software when the function call is evaluated.

True functional languages conform to these ideas both by what they do provide, such as pattern expressions, and what they do not provide, such as procedural iteration or dynamic assignments. The classic tutorial example of the factorial function, for example, could be expressed in the Haskell language by the pattern:

```
factorial x = if x > 0
              then x * factorial (x-1) else 1,
```

plus some type information, such as that a value for x must be an integer scalar.

Is R a functional programming language in this sense? No. The structure of the language does not enforce functionality; Section 2.3 examines that structure as it relates to functional programming and OOP. The evolution of R from earlier work in statistical computing also inevitably left portions of earlier pre-functional computations; Section 3 outlines the history. Random number generation, for example, is implemented in a distinctly “state-based” model in which an object in the global environment (`.Random.seed`) represents the current state of the generators. Purely functional languages have developed techniques for many of these computations, but rewriting R to eliminate its huge body of supporting software is not a practical prospect and would require replacing some very well-tested and well-analyzed computations (random number generation being a good example).

Functional programming remains an important paradigm for statistical computing in spite of these limitations. Statistical models for data, the motivating example for many features in S and R, illustrate the value of analyzing the software from a functional programming perspective. Software for fitting models to data remains one of the most active uses of R. The functional validity of such software is important both for theoretical justification and to defend the results in areas of controversy: Can we show that the fitted models are well-defined functions of the data, perhaps with other inputs to the model such as prior distributions considered as additional arguments? The structure of R as described in Section 2.3 can provide support for analyzing functional validity. Equally usefully, such analysis can also illuminate the limits of functional validity for particular software, such as that for model-fitting.

2.2 Object-Oriented Programming

The main ideas of object-oriented programming are also quite simple and intuitive:

1. Everything we compute with is an *object*, and objects should be structured to suit the goals of our computations.
2. For this, the key programming tool is a *class* definition saying that objects belonging to this class share structure defined by *properties* they all have, with the properties being themselves objects of some specified class.
3. A class can *inherit* from (contain) a simpler superclass, such that an object of this class is also an object of the superclass.

4. In order to compute with objects, we can define *methods* that are only used when objects are of certain classes.

Many programming languages reflect these ideas, either from their inception or by adding some or all of the ideas to an existing language.

Is R an OOP language? Not from its inception, but it has added important software reflecting the ideas. In fact, it has done so in at least three separate forms, giving rise to some confusion that this paper attempts to reduce.

Some of the confusion arises from not recognizing that the final item in the list above can be implemented in radically different ways, depending on the general paradigm of the programming language. A key distinction is whether the methods are to be embedded in some form of functional programming.

Traditionally, most languages adopting the OOP paradigm are not functional; either the language began with objects and classes as a central motivation (SIMULA, Java) or added the paradigm to an existing non-functional language (C++, Python). In such languages, methods were naturally associated with classes, essentially as callable properties of the objects. The language would then include syntax to call or *invoke* a method on a particular object, most often using the infix operator “.”. The class definition then encapsulates all the software for the class. Where methods are needed for other computations, such as special method names in Python or operator overloading in C++, these are provided by ad-hoc mechanisms in the language, but the method remains part of the class definition.

In a language that is functional or that aspires to behave functionally as S and R do, the natural role of methods corresponds to the intuitive meaning of “method”—a technique for computing the desired result of a function call. In functional OOP, the particular computational technique is chosen because one or more arguments are objects from recognized classes.

Methods in this situation belong to functions, not to classes; the functions are *generic*. In the simplest and most common case, referred to as a standard generic function in R, the function defines the formal arguments but otherwise consists of nothing but a table of the corresponding methods plus a command to select the method in the table that matches the classes of the arguments. The selected method is a function; the call to the generic is then evaluated as a call to the selected method.

We will refer to this form of object-oriented programming as *functional* OOP as opposed to the *encapsulated* form in which methods are part of the class definition.

2.3 Their Relationship to R

To understand computations in R, two slogans are helpful:

- Everything that exists is an object.
- Everything that happens is a function call.

In contrast to languages such as Java and C++ where objects are distinct from more primitive data types, every reference in R is to an object, in particular, to a single internal structure type in the underlying C implementation. This applies to data in the usual sense and also to all parts of the language itself, such as function definitions and function calls. Computations that are more complex than a constant or a simple name are all treated as function calls by the R evaluator, with control structures and operators simply alternative syntax hiding the function call. [Details and examples are shown in (Chambers, 2008, pages 458–468).]

The two slogans, however, do not imply that computations in R must follow either functional or object-oriented programming in the senses outlined in the preceding sections. With respect to object-oriented programming, R has several implementations that have evolved as outlined in Section 3. These can be used by programmers to provide software following either of the OOP paradigms.

Functional programming's relationship to R is less straightforward. The evaluation process in R does not enforce functional programming, but does encourage it to a degree. In particular, the evaluation process in R contributes to functional programming by largely avoiding side effects when function calls are evaluated, but some mechanisms in the language and especially in the underlying support code can behave in a non-functional way. To understand in a bit more detail, we need to examine this evaluation process.

Computations in R are carried out by the R evaluator by evaluating function call objects. These have an expression for the function definition (usually a reference to it by name) and zero or more expressions for the arguments to the call. The full details are somewhat beyond our scope here, but an essential question is how references to objects are handled. Any programming language must have references to data, which in R means references to objects. As discussed in Section 3, the evolution of such references is central to

the evolution of programming languages, especially for statistics.

In R a reference to an object is the combination of a name and a context in which to look up that name; the contexts in R are themselves objects, of type “environment”. A reference is therefore the combination of a name and an environment. (We'll look at an example shortly.)

Note that we are talking about references *to* objects; most objects in R are not themselves reference objects. Languages implementing OOP in the traditional, non-functional form essentially always include reference objects, in particular, what are termed *mutable* references. If a method alters an object, say, by assigning new values to some of its properties, all references to that object see the change, regardless of the context of the call to the method. Whether the reassignment of the property takes place where the object originated or down in some other method makes no difference; the object itself is the reference.

In contrast, the reference in R consists of a name and an environment—the environment in which the object referred to has been assigned with that name. Most R programming is based on a concept of *local references*; that is, reassigning part of an object referred to by name alters the object referred to by that name, but only in the local environment. If that local reference started out as a reference in some other environment, that other reference is still to the original object.

To understand the relation of local references to functional programming in R, an example and a few more details of function call evaluation are needed. R evaluates function calls as objects. For example, when the evaluator encounters the call

```
lm(Sepal.Width
   ~ . - Sepal.Length, iris),
```

it uses the object representing the call to create an environment for the evaluation.

The call identifies the function, also an object of course, typically referring to it by name. In this case `lm` refers to an object in the stats package. That object has formal arguments [14 of them, in the case of `lm()`]. The evaluator initializes an environment for the call with objects corresponding to the formal arguments, as unevaluated expressions built from the two actual arguments and default expressions found in the function definition. For details see Section 4 of the language definition, R Core Team (2013) and Chapter 13 of Chambers (2008). As an aside, the common use of terms like “call by value” (and the contrasting “call by

reference”) for argument passing in R is invalid and misleading. Arguments are not “passed” in the usual sense.

Local references operate on all the objects in the environment to prevent side effects. The formal argument `data` to `lm()` matches the expression `iris`, which refers to an object in the `datasets` package. Expressions that extract information from `data` work on that object. But the local reference defined by `data` and the environment of the evaluation is distinct from the reference to `iris` in the package. If an assignment or replacement expression is encountered that would alter `data`, the evaluator will duplicate the object first to ensure locality of the reference.

The local reference paradigm is helpful in validating the functionality of an R function. Only the local assignments and replacements need to be examined; calls to other functions will not alter references in this environment, so long as those functions stick to local reference behavior. If a function `f()` calls a function `g()` and both functions stick to local reference assignments, then knowing that the value of a call to `g()` depends only on the arguments is all that is needed; how `g()` computes that value is irrelevant.

While local references help avoid side effects, they do not prevent computations from referring to objects or other data outside the functions being called, and therefore potentially returning a result that depends on a non-functional “state.” Whether a particular computation in R is strictly functional can only be determined by examining it in detail, including all the functions that call code in C or Fortran.

The rest of this section takes a slight detour to consider how one might do that examination.

Validating Functionality in R

In principle, the functional validity of particular computations could be analyzed and either certified or the limitations to functionality reported. Such functional validation would be useful in cases where either the theoretical validity or the implications of the result in an application are being questioned. Fitting models to data provides a natural example for both aspects. Given a function taking as arguments `data` and a model specification and returning a fitted model object, can one validate that the returned object is functionally defined by the arguments? If not, can the non-functionality be parametrized meaningfully, in which case one can construct a functional version of the computation by including such parameters as implicit arguments? R does not have organized support for such

validity investigations, but developing tools for the purpose would be a worthwhile project.

Functional validation is a bottom-up construction. The bottom layer consists of any functions called that are not implemented in R, typically those that call routines in C++, C or Fortran. Included are the R primitives, routines from numerical libraries and a variety of other standard sources, plus any new code brought in to implement the computation in question. The functional validity of each of these is an empirical assertion. Some are clearly non-functional, such as the “`<-`” operator and `assign()` function that do nonlocal assignments.

Many computations in R eventually call subprograms not originally written for R. Each of these must be examined for potential non-functional behavior, sometimes a daunting task. However, good practice in using well-tested, preferably open-source supporting software will often provide a plausible basis.

If R code includes an interface to code in C, Fortran or other languages whose functional validity cannot be established, nothing more can be said. Other than such code, functional validity is likely to fail for one of three reasons:

- dependance on nonlocal values;
- using low-level computations in R known to violate functionality;
- changing functions or other objects at run time.

A prime example of the first is the use of external data, such as the global options object, for convergence tolerances or other parameters for iterative numerical computations. An example of the second is the inclusion of pseudo-random values in the calculation. The third problem might be caused, for example, by using a function from the global environment.

The third danger is greatly reduced when the code resides in the namespace of a package with explicit import rules. Any reasonable approach to validating functionality would make this a requirement.

My feeling is that most examples of failures could be corrected to create functionally valid extensions of the computation in question. Tolerances are often organized through the `Roptions()` function, explicitly designed to avoid functional programming by allowing users to set state parameters that are then queried by the calculation. Once identified, such options could be converted to additional arguments to the function being validated. [A general mechanism would be a version of `getOption()` that required the option in question to be supplied as an argument.]

Pseudo-random values are used in a variety of procedures, including some optimization techniques where they are expected to provide more robust numerical behavior by jittering values during iteration. These can be made functionally valid by using well-defined generator software, such as that supplied in R itself, and by treating the initial state of the generator as another nonlocal value to be incorporated as an additional argument. One should always include an explicit initialization via `set.seed()` in any example expected to be reproducible, and that practice can be the basis for a functionally valid version of the computation.

Beyond these specific examples, numerical computations often depend on the underlying parameters of the floating-point computations, for example, to select convergence criteria for iteration. Fortunately, several decades of work by numerical analysts and hardware designers have greatly standardized the specification of the numerical engine in modern computers: just knowing 32-bit or 64-bit gets us a long way.

Developing a framework for validating functionality seems to me an interesting cooperative research direction that could be of value to the statistical community.

3. THE EVOLUTION OF FUNCTIONAL PROGRAMMING, OOP AND R

The computational paradigms for functional programming and for object-oriented programming have evolved from a sequence of changes in software, beginning with the earliest programable computers. During the same period, software for statistics was also evolving, one thread of which led through early libraries to S and then to R.

There may be an appearance of earlier languages being replaced by later and presumably improved approaches. It is true that each major revision asserts improvements that will extend our abilities to express our ideas in software. However, none of the versions of S or R actually totally replaced earlier software paradigms.

The current software in, and interfaced from, R illustrates this evolution. R has developed important new techniques, but originated from the S language, reproducing nearly all of S as it was described at that time. S in turn went through several evolutionary changes and was itself based on extensive earlier software, particularly subroutine libraries for Fortran programming. Examining the history shows that a surprising portion of what we see now is structure inherited from the early stages.

The form in which functional programming and OOP were adopted was also influenced by the existing software. Examining the history will explain many of the choices made.

3.1 From Hardware to Data and Libraries

The earliest general-purpose computers were programmed in terms of the physical machine, its storage and the basic operations provided to move data around and perform arithmetic and other operations. The IBM 650 (Figure 1) was probably the first computer widely sold and used (and the machine on which I did my first programming, around 1960).

In this pre-silicon world, storage for data or programs resided on a rotating magnetic drum, holding 2000 decimal words. Data could be read or written only when the corresponding segment of the drum passed under the appropriate fixed head, so that physical positioning of data was a serious aspect of performance. With this close view of the hardware, programming languages (assembly languages for the actual machine instructions) defined storage in terms of single physical units (words in the 650) and blocks of sequential storage.

This was not an environment to encourage abstraction of ideas about data. However, by 1960 the first generation of “high-level” languages had been introduced and would support profound changes. For statistical computing this meant primarily Fortran.

In terms of data storage, Fortran actually continued the basic notion of single items (scalars) and contiguous blocks (arrays). Two major changes, however, were made. First, the contents were described in terms of their content, the first *data types* including integer and floating point numbers. Second, the language encouraged operations that iterated over the contents of the arrays. By interpreting an array as a sequence of equal-length subarrays, this indexing extended to matrices and to multi-way tables.

Along with the new paradigm for data and facilities for iteration, the high-level languages encouraged software to be organized in subroutines, so that a computational method could be realized as one or several units of software. While the changes may seem modest from the current perspective, they in fact supported a major revolution in scientific computing generally and emphatically so in computing for statistics.

Algorithm series and other publications supported by professional societies began to accumulate refereed, trustworthy procedures for many key computations. The statistics research group at Bell Labs developed a large Fortran library that reflected our needs



FIG. 1. An IBM 650 computer, mid 1950s. Under the glass is the magnetic drum storage unit (memory), 2000 words for data and programs.

and our philosophy of research and data analysis. The book “Computational Methods for Data Analysis”, Chambers (1977), did not present software but did reflect the tools that would later form the basis for S. After an introduction and discussion of program design, the remaining six chapters covered computations supported by the library:

3. *Data Management and Manipulation* (including sorting and table lookup).
4. *Numerical Computations* (approximations, Fourier transforms, integration).
5. *Linear Models* (numerical linear algebra, regression, multivariate methods).
6. *Nonlinear Models* (optimization, nonlinear least squares).
7. *Simulation of Random Processes* (random number generation and Monte Carlo).
8. *Computational Graphics* (plotting techniques, scatter plots, histograms and probability plots).

Each of these was supported in the pre-S era by subroutines that would then become the basis for corresponding functions in S.

Much of the organization for basic tools in R has inherited, through S, the structure of the subroutine

library. That includes the graphical computations, in particular, features essential to S and R: separation of graphic device specification from plotting; the plot, figure and margins structure; graphical parameter specification to control style. These were not created for S but taken over from previous Fortran software, described in Becker and Chambers (1977).

The Bell Labs software was in the background of Chambers (1977), but general readers were given instructions for obtaining similar software from publicly available sources for the methods described. The procedure would not always be simple, but the potential availability marked a big step forward. For the first time, statisticians could draw on an extensive range of relevant software to support their research, at least in principle. Various statistical software packages had existed for some time, but these were by and large oriented to routine analysis, to teaching or to specialized statistical techniques. Chambers (1977) and the software it reflected were aimed at research in statistics and challenging data analysis. For this purpose, a more general and open-ended approach was needed.

3.2 From Fortran to S

For those involved with statistical theory or applications, in academia or industry, there were two main

limitations to the software described so far: availability and the programming interface. The Appendix to Chambers (1977) was a set of tables for each of the chapters, with rows corresponding to computational tools that were more or less available to readers. The last column of the table listed sources for the corresponding software. The entries in that column were not uniformly helpful; in the best situation, a generally available program library could be ordered that provided a number of the subroutines, but these were not designed for statistical applications, most being directed at numerical methods typically motivated by applications in physics. More than half of the entries read "Listing," implying a laborious and error-prone manual procedure for the user. [As an example, many "bug reports" came to us as a result of confusing an "I" and a "1" when typing in the stable distribution software, Chambers, Mallows and Stuck (1976).]

Substantial in-house libraries, such as the one at Bell Labs, gave users a fairly wide range of computations, supported by improved numerical and other algorithms. However, to apply the computations specifically to a particular dataset with particular results in mind required some substantial additional Fortran programming. That programming had to be repeated and revised for each analysis or research question.

In the 1970s the situation was therefore a combination of improved basic computational capabilities but with a high programming barrier for most statisticians. The classical linear regression in Fortran as shown in Becker and Chambers (1985), for example, was fairly straightforward:

```
call lsfit(X, N, P, y, coef, resid).
```

This computes the fitted model and returns it as vectors of coefficients and residuals. The data as objects are restricted to arrays, a matrix *X* and vector *y* for the data and two arrays, *coef* and *resid* for the fitted model. The structure of the objects and their storage allocation remains the programmer's responsibility. Linking the basic computation to the data in an actual analysis remained nontrivial and mistakes along the way were likely. And this is for the most standard of models. Even given an extensive library, the programming to apply the tools to most applications was a laborious, error-prone activity, usually assigned to dedicated programmers, research assistants or students. The statistician's ideas went through nontrivial translation before they were expressed as computations.

The first two versions of *S* were designed to provide an "interactive environment" that included the computational areas described in Chambers (1977) and that

allowed the statistician to formulate ideas directly for computation. The second version of *S* was licensed for general use and described in Becker and Chambers (1984).

In *S*, the linear regression computation became a simpler expression, storage for data was provided automatically and the returned model was now an object, with components for the coefficients and residuals:

```
fit <- reg(X, y).
```

At this stage, *S* had a functional appearance, not radically unlike *R*, but its paradigm was essentially an extension of the Fortran view. Dynamically created, self-describing objects were assigned in a single workspace, but the underlying computations were those of the earlier subroutine library: The functions in *S*, documented in Becker and Chambers (1984), were in fact interfaces to Fortran subroutines: `reg()` would in fact be programmed by calling `lsfit()`.

Although there was a macro facility in the language, programming a function in this version of *S* meant "extending *S*" as described in the book of that name, Becker and Chambers (1985). The definition of the new function was programmed in an "interface language" built on Fortran and compiled from its Fortran translation. As the main programming mechanism this was unsatisfactory, in the sense that extending the language had a substantial learning barrier beyond using the language. The ability to access other software via an inter-system interface remains a key feature of *R*, however, one still under active development.

Equally as important as the technical side was the beginning of a network of statisticians involved in creating and sharing software through the medium of the language. *S* was licensed from the early 1980s, available thanks to the newly distributed UNIX operating system, with inexpensive academic licenses to encourage adoption by university researchers, also following the example of UNIX. Open-source software was not an option, but the research community was increasingly involved and their interest stimulated further developments on our part, particularly from contacts with interested users belonging to a "beta testing" network.

Simultaneously, we were thinking about a new approach to the language itself, emphasizing the programming aspect of creating new software for statistical and other quantitative applications. Described initially in Chambers (1987) as a language separate from *S*, this research later merged with other changes to form the next version, labeled *S3* and described in the "blue book," Becker, Chambers and Wilks (1988).

The slogans in Section 2.3 were basic to this version of S: everything is an object (stated explicitly) and function calls do all the computation (implicit).

This was functional programming (more or less) and object-*based* but not object-oriented. Objects were given structure through attributes attached to vectors and through named components, but there were no classes or methods.

3.3 From Data to Classes and Methods

The languages that originated the concepts of classes, properties, inheritance and methods came out of several motivations. The first, Simula, was concerned with simulating systems. In retrospect, modeling by simulation and modeling by fitting to data have clear correspondences but with quite a different perspective. For an example, suppose we want to simulate a simple model for an evolving population of individuals. In R notation, but quite in the style of Simula, we define a class `SimplePop`. An object from this class is a specific realization of the model population with properties that define the probabilities of birth and death, and a vector of population size at each generation. An object from the population is created by calling the generator for the class:

```
p <- SimplePop(birth = 0.08,
               death = 0.1,
               size = 100).
```

Rather than a single functional computation as in the case of linear regression, computations proceed by simulating the evolution of the population object `p`. The object itself evolves; in the terminology of OOP, it is a *mutable reference*.

A corresponding difference in the programming paradigms of S and the emerging OOP languages was that the latter did not take a functional view of computation. Instead, computations largely consisted of invoking a method on an object. In the `SimplePop` example, the fundamental computation is to simulate one generation of the evolution by invoking the `evolve()` method

```
p$evolve().
```

The value returned by this method is irrelevant. The method's purpose is to change the object, in this case by simulating one further generation and appending the resulting value to a property in the object, namely, `p$size`. (See files "`SimplePop.R`" and "`SimplePopExample.R`" in the supplementary materials.)

Following the development of Simula in the late 1960s, a variety of languages adopted this paradigm. C++ added classes and methods to the C language; like C, it was initially used for a variety of programming tasks implementing UNIX and application software for UNIX. In contrast to the "add-on" nature of C++, the Smalltalk language was a very pure, simplified realization of the ideas in Simula. Its major, and revolutionary, application was to implement the graphical user interface created at Xerox PARC in the 1970s. Many other versions of encapsulated OOP followed, either added on to existing languages or incorporated into new languages from the start.

Dialects of the Lisp language and languages based on Lisp also incorporated OOP in various forms. During the 1980s, several research projects built statistical software on the basis of these languages, including some elegant and potentially widely applicable systems, notably LISP-STAT, Tierney (1990). As it turned out, however, the most widely used version of OOP for statistical applications would come from a somewhat casual approach in S.

3.4 Functional OOP in S and R

The chief motivation for introducing classes and functional methods to S was the initial application: fitting, examining and modifying diverse kinds of statistical models for data. This remains arguably the most compelling example for functional OOP in statistics. The "Statistical Models in S" project reported in Chambers and Hastie (1992)—the "white book"—brought together ten authors presenting software for a variety of statistical models, from linear regression to tree-based models. The different models were presented as consistently as possible.

Each type of model had a definition as an object having the information, such as coefficients and other properties, required. The object was created by a corresponding function taking as arguments the data, model description and possibly other controlling parameters. A linear regression fit, for example, called the function `lm()`:

```
irisFit <- lm(Sepal.Width
             ~ . - Sepal.Length, iris)
```

and returned a corresponding linear regression object. Further computations on this object would examine the model, return information about it, or update the fit. The underlying computations still used basic software similar to that for `lsfit()` and `reg()`. However, the description of the model (a formula) and the

data (a data frame) were designed to apply to statistical models generally. For example, to fit a generalized linear model the user called `glm()` with formula and data arguments typically similar to those in a call to `lm()`. Other arguments would provide information suitable to the particular type of model (a `link` function, e.g.).

For the convenience of the user, further computations should have a uniform appearance. To print or plot the fitted model or to compute predictions or an updated model corresponding to new data, the user should call the same function [`print()`, `plot()`, `predict()` or `update()`] in the same way, regardless of the type of model. The owner of the software for a particular type of model, on the other hand, would like to write just that version of each function, without being responsible for the other versions.

Once stated, this is essentially a prescription for functional OOP: a class of objects for each kind of model, generic functions for the computations on the objects and methods for each function for each class. Where one class of models is an extension of another (analysis of variance as a subclass of linear models, e.g.), methods can be inherited when that makes sense.

An implementation of generic functions and methods was introduced as part of the statistical models project and described in the Appendix to the white book. The central mechanism was an explicit method dispatch. The function `print()`, for example, would evaluate the expression:

```
UseMethod("print").
```

The evaluation of this call would examine the “class” attribute of the first formal argument to the function. If present, this would be a character vector. Eligible methods would be those matching one of the strings in the class vector; if none matched, a method matching the string “default” would be used. Inheritance was implemented by having more than one string in the class, with the first string being “the” class and the remainder corresponding to inherited behavior.

Chambers and Hastie (1992), in the discussion of classes and methods, noted that S differed from other OOP languages because of its functional programming style. In fact, this version of functional OOP finessed the resulting distinction from encapsulated OOP in two ways. First, the methods were dispatched according to a single argument, the first formal argument of the generic function in principle. As a result, the methods were unambiguously associated with a single class, as they would be in encapsulated OOP. Methods were actually dispatched on either argument to the usual bi-

nary operators, but a number of encapsulated OOP languages do the same, under the euphemism of operator overloading.

Second, the question of whether methods belonged to a class or a function was avoided by not having them belong to either. Methods were assigned as ordinary functions and identified by the pattern of their name: “*function.class*”. In any case, there were no class objects and generic functions were ordinary functions that invoked `UseMethod()` to select and call the appropriate method. Neither the function nor the class was able to own the methods.

Technically, the method dispatch in this version of OOP was instance-based, not class-based, since no rule enforced a consistent set of classes, that is, that all objects with a given first class string would have identical following strings for the superclasses. (R for some time had an S3 class in the base package with a main class string “POSIXt”, representing date/times, that could be followed in different objects by one of two strings that in fact represented specializations, i.e., subclasses, of “POSIXt”.)

The classes and methods implemented for statistical models constituted a bare-bones version of functional OOP, which is not to imply that this was a bad idea. Advantages include a relatively low learning barrier for programming and a thin implementation layer above the previously existing language, which in turn means less computational overhead in some circumstances. [Interestingly, the encapsulated OOP of Python has a similarly thin implementation, with classes containing methods but without defining the properties. A very analogous defense is made for that implementation, in Section 9 of the Python tutorial, Python (2013), e.g.]

A more formal version of functional OOP was developed at Bell Labs, introduced into S in the late 1990s and described in Chambers (1998). By this time, S-based software was exclusively licensed to the Insightful Corporation, which later purchased the rights to the S software, in 2004, and was itself subsequently purchased by Tibco.

The new paradigm differed from S3 classes and methods in three main ways:

1. Methods could be specified for an arbitrary subset of the formal arguments, and method dispatch would find the best match to the classes of the corresponding arguments in a call to the generic function.
2. Classes were defined explicitly with given properties (the slots) and optional superclasses for inheriting both properties and methods.

3. Generic functions, methods and class definitions were themselves objects of formally defined classes, giving the paradigm reflectivity.

The new paradigm was part of the version of S described in the 1998 book and generally referred to as S4. The S4 label is generally applied to this OOP paradigm, whether in S or R. S4 methods never had much chance of replacing S3 methods. In practice, many S4 generic functions were based on functions that already dispatched S3 methods. In this case, the S3 generic function became the default S4 method.

The work on S4 paralleled in time the arrival of R and its conversion into a broad-based joint project following the initial publication by Ihaka and Gentleman (1996). The implementation of R was designed to provide the functionality for S described in the blue book and white book, including S3 methods. Beginning in 2000, an implementation of the S4 version of OOP was added to R. The “Software for Data Analysis” book, Chambers (2008), includes a description of the R version.

Both versions of functional OOP will remain in R. Many prefer the simplicity of the old form, and in any case the very large body of existing code will not be discarded, and should not be. Some important extensions have been made, for example, by registering the S3 methods from a package. Major forward-looking projects have typically used the newer version, for example, the Bioconductor project for bioinformatics software, Gentleman et al. (2004), and the Rcpp interface to C++, Eddelbuettel and François (2011). Recent changes, such as making the S3 and S4 versions of inheritance as compatible as possible, have been aimed at helping the two forms to coexist productively.

Any programming paradigm with some degree of formality is likely to have a higher initial learning barrier and require some extra specification from the programmer. A comparison of encapsulated OOP programming with Python to that with Java is an interesting parallel to S3 and S4. In both examples, the less formal version is likely to be quicker to learn, while the more formal version provides more information about the resulting software. That information in turn can support some forms of validation for the resulting software, as well as tools to analyze and describe it. Python and Java being rather different languages in other respects as well, projects are not too likely to make a choice between them based solely on the formality of the object-oriented programming.

With R, a conscious choice is more likely. The arguments for a more formal approach apply particularly, in

my opinion, to projects with one or more of the characteristics: a substantial amount of software is likely to be written; the application has a fairly wide scope in terms of either the data or the computing methods; or the validity and reliability of the resulting software is important.

Nothing prevents good software being written without formal tools in this case nor of bad software being written with them. However, there are several potential benefits that can be summarized in parallel with the main innovations noted above:

1. Allowing methods to depend on multiple arguments fits the functional paradigm in R, in which the arguments collectively define the domain of the function. Many functions in R are naturally applied to different classes of objects, not necessarily corresponding to the first argument, or only to one argument. For example, when binary operators such as arithmetic are defined for a new class, a clean design of methods for the operators often needs to distinguish three cases: the first operand only belonging to the new class, the second operand only or both operands.

2. A formal definition for a class allows programmers to rely on the properties of objects generated from the class. Otherwise, the nature of the objects can only be inferred, if at all, from analyzing all the software that creates or modifies an object of this class.

3. Having formal definitions for the generic functions, methods and class definitions themselves supports a growing set of tools for installing and using packages that include such functions, methods or classes.

The benefits of a general, reliable form of functional OOP extend to developments in the language itself. For example, reference classes were built on the S4 classes and methods, with no internal changes to the R evaluator required.

3.5 Reference Classes

Functional OOP remains an active area in R. In addition, *reference classes*, introduced to R in 2010 in version 2.12.0, provide an implementation of encapsulated OOP. Class definitions include the properties of the class with optional type declarations; properties may also be optionally declared read-only. Class definitions are themselves objects available at runtime. Methods are programmed as R functions, in which the object itself is implicitly available, not an explicit argument. Methods can access or assign properties in the

object by name. These characteristics make the implementation more Java-like, say, than Python- or C++-like.

The programmer defines a reference class in the R style, calling `setRefClass()` instead of `setClass()`. The call returns a generator for the class and saves the class definition object as a side effect, as does `setClass()` for S4 classes.

As a side comment, while R uses a model for most of its objects and computations that is fundamentally different from the object references in encapsulated OOP, a few key features made the implementation of reference classes in R possible and even relatively straightforward. Most importantly, the R data type “environment” provides a vehicle for object references and properties. Environments are universal in R and well supported by programming tools. In particular, the active binding mechanism, which allows access and assignment operations on objects in environments to be programmed in R, was valuable in the implementation.

Reference classes allow the use of encapsulated OOP for objects that suit that paradigm more naturally than they do functional OOP. As noted in Section 3.3, the essential distinction between functional and encapsulated OOP is whether an object is created, once, by a function call or is instead a mutable object that changes as methods are invoked.

Statistical computing has examples clearly suited to each of these paradigms. The linear model returned by `lm()` is not open to mutation. Change the numbers in the coefficients or residuals and you no longer have an object that should belong to that class. In contrast, a model simulating a dynamic process such as the `SimplePop` class in Section 3.3 exists precisely for the purpose of changing, with its evolution being the central point of interest. Other, less directly statistical computations in R also may correspond to mutable objects, for example, the frames or other objects in a graphical interface.

Not every case is clear cut. Sometimes, essentially the same class structure may be more appropriate for functional or encapsulated classes depending on the purpose of the computation. Data frames are a prime example. This essential object structure is viewed naturally as functional when it is part of a functional object related to the data frame. For example, a fitted model that wanted to be fully reproducible could return the data frame on which the fitting was based [e.g., `lm()` includes the model frame it constructs]. Such a data frame is clearly functional; again, change it and you

invalidate the model. On the other hand, a data frame to be used in data cleaning and editing is an object that needs to be mutable.

Having both paradigms in a single language is unusual. Some functional-style languages have implemented functional OOP, notably Dylan, interesting for its parallels with OOP in R—see Shalit (1996), particularly the discussion of method dispatch. Other languages with a functional structure have nevertheless added what is essentially encapsulated OOP, for example, Odersky, Spoon and Venners (2010) for the case of Scala.

We hope that providing both paradigms in R encourages software design that is natural for the application. It does at the same time pose some subtleties. Reference classes and reference class objects are somewhat abnormal in R. One needs to understand the distinctions from standard R objects.

The key is the local reference mechanism noted in Section 2.3. The R evaluator enforces local reference by duplicating an object when a computation might alter a nonlocal reference. Certain object types are exceptions that are not duplicated. The important exception is type “environment”. Reference classes are implemented by extending this type. Encapsulated OOP in R uses no special form of the function call. Method invocation is just a call to the “\$” operator, for which reference classes have an S4 method. Reference semantics are obtained by one basic fact: environments are never duplicated automatically. The S4 class mechanism in R nevertheless allows one to subclass the “environment” type in order to define reference class behavior.

The objects in the fields of a reference class object can be ordinary R objects. They behave just as usual and when used in function calls will have regular local reference behavior in that call. It is only when fields in the reference object itself are replaced that the encapsulated OOP is relevant.

Reference class objects are also good candidates for interfaces to other languages that implement the same OOP paradigm, such as Java, C++ or Python. The R object could be a proxy for an object in the other language with methods invoked in R but executed on the original object. The Rcpp interface to C++, Eddelbuettel and François (2011), has a mechanism for extending C++ classes in this way. C++ classes can only be inferred from the source, meaning that either the programmer must supply the interface information (as in the current implementation) or some processing of the source must be applied (currently used to export

functions from C++ but not classes). Java classes are accessible as objects, via “reflectance” in Java terminology, so that in principle proxy classes in R should be possible. The rJavax package by Danenberg (2011) has an initial implementation. For Python, methods are available from the objects but properties are not formally defined. At the time of writing, basic interfaces to Python exist, for example, Grothendieck and Bellostta (2012), which could be extended to support class interfaces, with methods but not properties inferred from the Python class objects.

Further work on these and other inter-system interfaces would be a valuable contribution to the user community.

4. SUMMARY

R plays a major role in the communication and dissemination of new techniques for statistics and for results of statistical research more generally. In particular, the many packages written in R or using R as a base for interfacing to other software constitute an essential, rapidly growing resource. Therefore, the quality of such software and the ability of programmers to create and extend it are important.

The current R language and its supporting functionality are the result of many years of evolution, from early programming libraries through the S language to R, which itself has evolved and accumulated a variety of programming techniques. This evolution has been much influenced by the functional and object-oriented programming paradigms. New versions have continued to include supporting software and programming tools found useful at earlier stages along with improved capabilities.

The programming paradigms become especially relevant when the applications are complex or the quality of the resulting software is important. In particular, the versions of object-oriented programming in R can assist in dealing with complexity of the underlying data. As noted, R implements OOP in two forms, functional and encapsulated. These are complementary, with one or the other suitable for particular applications. The latter is essentially the form of OOP used in most other languages, but the former is distinctly different. Considerable confusion has arisen in discussions of OOP in R from not noting that distinction, which the present paper has tried to clarify.

More generally, understanding the role of object-oriented and functional programming in R may assist

future contributing programmers in using related programming tools. The continuing rapid growth of R-based software and the expanding, challenging range of techniques it has to support make effective programming an important goal for the statistical community.

The importance of object-oriented programming is likely to increase as statistical software takes on new and challenging applications. In particular, the need to deal with increasingly large objects and distributed sources of data will bring in specialized classes of data and will need powerful computing tools. One important direction has been to transform selected software in R, particularly to speed up large-scale computations; see, for example, the companion paper Temple Lang (2014). Complementary to this is to interface to other languages and software when these provide better performance on “big data” and other computationally demanding applications. In particular, interfaces that match with object-oriented treatments for specialized forms of data can exploit the OOP facilities in R. The interface to C++, Eddelbuettel and François (2011), is an example. Further development of such interfaces will be of much benefit.

Functional programming is perhaps not such an obviously hot topic at the moment. However, the underlying philosophy that our software should be in the form of reliable, defensible units is very much part of R. Situations where the validity of statistical computations needs to be defended are likely to increase, given the growing need for statistical treatment of complex problems for science and society.

ACKNOWLEDGMENTS

Thanks to the Associate Editor and the referees for some helpful comments on presentation and content. Thanks especially to Vincent Carey for organizing and encouraging the set of talks and papers of which this is part.

REFERENCES

- BECKER, R. A. and CHAMBERS, J. M. (1977). Gr-z: A system of graphical subroutines for data analysis. In *Proc. Interface Symp. on Statistics and Computing* 10 409–415.
- BECKER, R. A. and CHAMBERS, J. M. (1984). *S: An Interactive Environment for Data Analysis and Graphics*. Wadsworth, Belmont, CA.
- BECKER, R. A. and CHAMBERS, J. M. (1985). *Extending the S System*. Wadsworth, Belmont, CA.
- BECKER, R. A., CHAMBERS, J. M. and WILKS, A. R. (1988). *The New S Language*. Chapman & Hall, Boca Raton, FL.
- CHAMBERS, J. M. (1977). *Computational Methods for Data Analysis*. Wiley, New York. MR0659716

- CHAMBERS, J. M. (1987). Interface for a quantitative programming environment. In *Comp. Sci. and Stat., Proc. 19th Symp. on the Interface* 280–286.
- CHAMBERS, J. M. (1998). *Programming with Data: A Guide to the S Language*. Springer, New York.
- CHAMBERS, J. M. (2008). *Software for Data Analysis: Programming with R*. Springer, New York.
- CHAMBERS, J. M. and HASTIE, T., eds. (1992). *Statistical Models in S*. Chapman & Hall, Boca Raton, FL.
- CHAMBERS, J. M., MALLOWS, C. L. and STUCK, B. W. (1976). A method for simulating stable random variables. *J. Amer. Statist. Assoc.* **71** 340–344. MR0415982
- DANENBERG, P. (2011). rJavax: rJava extensions. R package version 0.3. Available at <http://CRAN.R-project.org/package=rJavax>.
- EDDELBUEITTEL, D. and FRANÇOIS, R. (2011). Rcpp: Seamless R and C++ integration. *Journal of Statistical Software* **40** 1–18.
- GENTLEMAN, R. C., CAREY, V. J., BATES, D. M. et al. (2004). Bioconductor: Open software development for computational biology and bioinformatics. *Genome Biology* **5** R80.
- GROTHENDIECK, G. and BELLOSTA, C. J. G. (2012). rJython: R interface to Python via Jython. R package version 0.0-4. Available at <http://CRAN.R-project.org/package=rJython>.
- IHAKA, R. and GENTLEMAN, R. (1996). R: A language for data analysis and graphics. *J. Comput. Graph. Statist.* **5** 299–314.
- ODERSKY, M., SPOON, L. and VENNERS, B. (2010). *Programming in Scala*, 2nd ed. Artima, Walnut Creek, CA.
- PYTHON (2013). The Python Tutorial. Python. Available at <http://docs.python.org/tutorial>.
- R CORE TEAM (2013). *R Language Definition*. R Foundation for Statistical Computing, Vienna, Austria. ISBN 3-900051-13-5. Available at <http://cran.r-project.org/doc/manuals/R-lang.html/>.
- SHALIT, A. (1996). *The Dylan Reference Manual*. Addison-Wesley, Reading, MA.
- TEMPLE LANG, D. (2014). Enhancing R with advanced compilation tools and methods. *Statist. Sci.* **29** 181–200.
- TIERNEY, L. (1990). *LISP-STAT: An Object-Oriented Environment for Statistical Computing and Dynamic Graphics*. Wiley, New York.